

Services

a Service is an abstraction which defines a logical set of Pods and a policy by which to access them. Services are using a selector field to target the Pods. Once the selector field of a service matches with labels assigned to the pods, these pods become the endpoints of the service.

Types of services:

1. NodePort
2. ClusterIP
3. LoadBalancer

Objectives:

1. Creating a deployment and exposing it to a NodePort service
2. Creating a deployment and exposing it to a ClusterIP service
3. Creating a deployment and exposing it to a LoadBalancer service

1. Creating a deployment and exposing it to a NodePort service:

It exposes the Service on each Node's IP at a static port. In other word, the traffic from outside world will use NodePort service to connect with the endpoints. Kubernetes allocates 30000-32767 ports to NodePort service

Step1: Create a deployment

We are going to create a deployment named **prod-deploy** with nginx image. Use the below command to create this deployment.

```
kubectl create deploy prod-deploy --image nginx --replicas 3
```

```
root@master:~# kubectl create deploy prod-deploy --image nginx --replicas 3
deployment.apps/prod-deploy created
root@master:~#
root@master:~# kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
prod-deploy   3/3     3            3           41s
root@master:~#
root@master:~# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
prod-deploy-5b9d84c685-9lhbq       1/1     Running   0           45s
prod-deploy-5b9d84c685-wcwsn       1/1     Running   0           45s
prod-deploy-5b9d84c685-zw94q       1/1     Running   0           45s
```

From above output, we can see that our deployment with 3 pods has been created. Now we will describe the deployment to check the labels which will be used in our service under selector field.

```
kubectl describe deploy prod-deploy
```

```
root@master:~# kubectl describe deploy prod-deploy
Name:                prod-deploy
Namespace:           default
CreationTimestamp:    Fri, 13 Jan 2023 06:19:45 +0000
Labels:              app=prod-deploy
Annotations:         deployment.kubernetes.io/revision: 1
Selector:            app=prod-deploy
Replicas:            3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:        RollingUpdate
MinReadySeconds:     0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=prod-deploy
  Containers:
    nginx:
      Image:      nginx
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available       True    MinimumReplicasAvailable
    Progressing     True    NewReplicaSetAvailable
OldReplicaSets:    <none>
NewReplicaSet:     prod-deploy-5b9d84c685 (3/3 replicas created)
Events:
  Type           Reason             Age   From                      Message
  ----           -
  Normal         ScalingReplicaSet  2m7s  deployment-controller     Scaled up replica set prod-deploy-5b9d84c685 to 3
```

From above output, we can see that the pods have the label **app=prod-deploy**

Step2: Create a NodePort service

We can simply use an imperative command or we can create a Yaml file to create a NodePort service.

```
kubectl expose deploy prod-deploy --name prod-service --type NodePort --port 80
--target-port=80
```

Above command will automatically expose our deployment and Kubernetes will assign the NodePort automatically from range 30000-32767.

Or Create the Yaml file

```
kubectl expose deploy prod-deploy --name prod-service --type NodePort --port 80
--target-port=80 --dry-run=client -o yaml > prod-service.yaml
```

```

apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: prod-deploy
  name: prod-service
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
      nodePort: 30010
  selector:
    app: prod-deploy
  type: NodePort
status:
  loadBalancer: {}

```

We have our Yaml file ready and let's understand the important fields inside this Yaml file.

Port: the port of the service

targetPort: port of the pod this service is targeting to.

NodePort: Port on the node which is exposed outside. We are explicitly using 30010, if we do not mention then Kubernetes will automatically assign the port.

Selector: It is used to target the Pods which are having the same label. So the pods having app=prod-deploy label they will become the endpoint of this service.

Now let's apply this Yaml file and we will check the status further

```
kubectl apply -f prod-service.yaml
```

Check the services available in our cluster in current namespace.

```
kubectl get svc
```

```

root@master:~# kubectl apply -f prod-service.yaml
service/prod-service created
root@master:~# kubectl get svc

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4d20h
prod-service	NodePort	10.97.174.141	<none>	80:30010/TCP	6s

```

root@master:~#

```

Our NodePort prod-service has been created and with IP address 10.97.174.141 which is provided by Kubernetes only. We can also see that 30010 nodePort has been assigned to this service.

Let's describe this service.

```
kubectl describe svc prod-service
```

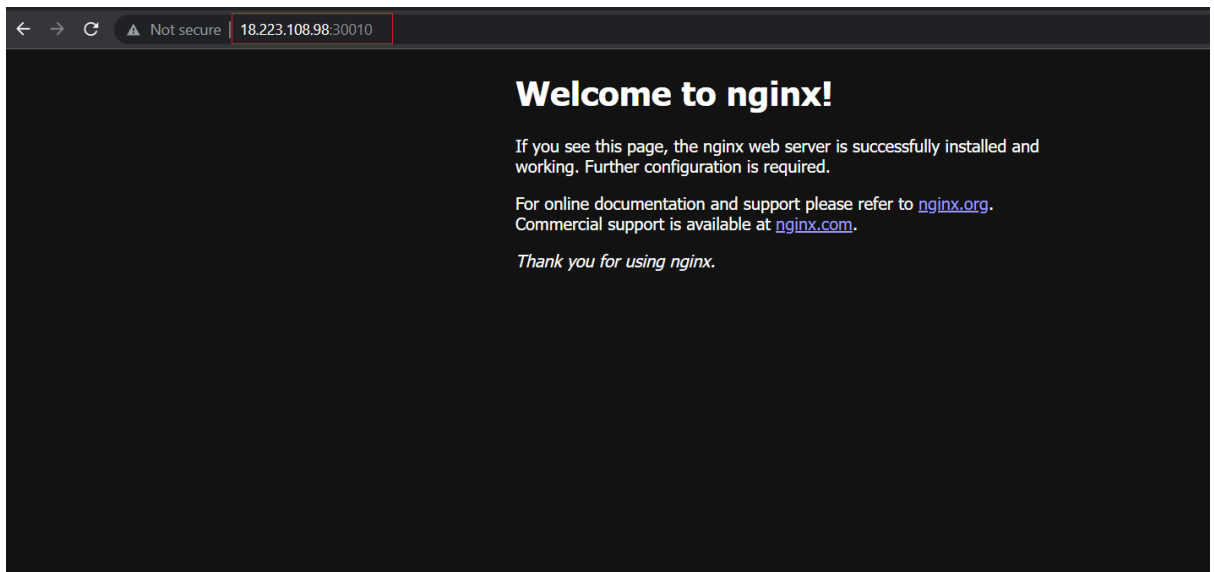
```
root@master:~# kubectl describe svc prod-service
Name:                prod-service
Namespace:            default
Labels:               app=prod-deploy
Annotations:          <none>
Selector:             app=prod-deploy
Type:                 NodePort
IP Family Policy:     SingleStack
IP Families:          IPv4
IP:                   10.97.174.141
IPs:                  10.97.174.141
Port:                 <unset> 80/TCP
TargetPort:           80/TCP
NodePort:             <unset> 30010/TCP
Endpoints:            192.168.189.74:80,192.168.235.129:80,192.168.235.134:80
Session Affinity:     None
External Traffic Policy: Cluster
Events:               <none>
root@master:~#
root@master:~#
root@master:~# kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
prod-deploy-5b9d84c685-9lhbq	1/1	Running	0	30m	192.168.235.134	worker1	<none>	<none>
prod-deploy-5b9d84c685-wcwsn	1/1	Running	0	30m	192.168.189.74	worker2	<none>	<none>
prod-deploy-5b9d84c685-zw94q	1/1	Running	0	30m	192.168.235.129	worker1	<none>	<none>

Above we can see that under the endpoints, we are having the same pods which are having **prod-deploy** label.

Now this has been exposed to the outside world. We can access it using the below in the browser.

node-ip-address:nodeport



So, we are able to access our application using NodePort service.

Now delete the service using below command.

```
kubectl delete svc prod-service
```

2. Creating a deployment and exposing it to a ClusterIP service

It exposes the Service on a cluster-internal IP. It is used access from within the cluster. This is the default that is used if you don't explicitly specify a type for a Service. It uses the same selector concept to access the

Step1: Create a deployment

Create a deployment using below command.

```
kubectl create deploy backend-deploy --image nginx --replicas=3
```

It will create a deployment **backend-deploy** having 3 pods with **nginx** image. Let's describe the deployment and we will check use the same labels for our **clusterIP** service.

```

root@master:~#
root@master:~# kubectl get deploy
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
backend-deploy      3/3     3             3           15s
root@master:~#
root@master:~# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
backend-deploy-5489dcddc6-48c74    1/1     Running   0           21s
backend-deploy-5489dcddc6-dggnx    1/1     Running   0           21s
backend-deploy-5489dcddc6-jvdpj    1/1     Running   0           21s

```

```

root@master:~# kubectl describe deploy backend-deploy
Name:                backend-deploy
Namespace:            default
CreationTimestamp:    Fri, 13 Jan 2023 07:12:44 +0000
Labels:               app=backend-deploy
Annotations:          deployment.kubernetes.io/revision: 1
Selector:             app=backend-deploy
Replicas:             3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:         RollingUpdate
MinReadySeconds:      0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=backend-deploy
  Containers:
    nginx:
      Image:      nginx
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type           Status    Reason
    ----           -
    Available       True      MinimumReplicasAvailable
    Progressing     True      NewReplicaSetAvailable
OldReplicaSets:    <none>
NewReplicaSet:     backend-deploy-5489dcddc6 (3/3 replicas created)
Events:
  Type           Reason             Age           From              Message
  ----           -
  Normal         ScalingReplicaSet  2m40s        deployment-controller  Scaled up replica set backend-deploy-5489dcddc6 to 3

```

Above we can see that pods are having **app=backend-deploy** label which we will be used under **selector** field of our **ClusterIP** service.

Step2: Create a ClusterIP service

We can use the imperative command as well as the Yaml file to create the ClusterIP service.

Use the below command to simply expose backend-deploy.

```
kubectl expose deploy backend-deploy --name backend-service --port 80
```

Or we can create a Yaml template using below command.

```
kubectl expose deploy backend-deploy --name backend-service --port 80 --dry-run=client -o yaml > backend-service.yaml
```

The above command will create a template. We can see below that this service is having the selector field **app: backend-deploy** which will target the pod having the same label.

As we have discussed in NodePort, port is for the service itself and targetPort is the pod's port.

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: backend-deploy
  name: backend-service
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: backend-deploy
status:
  loadBalancer: {}
```

Let's apply this file and describe it to get more details.

```
kubectl apply -f backend-service.yaml
```

```
root@master:~# kubectl apply -f backend-service.yaml
service/backend-service created
root@master:~#
root@master:~# kubectl get svc
NAME                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
backend-service     ClusterIP   10.109.215.53   <none>           80/TCP           4s
kubernetes           ClusterIP   10.96.0.1       <none>           443/TCP          4d21h
root@master:~#
```

Above we can see that our clusterIP service has been created with **10.109.215.53** IP address which has been assigned by Kubernetes only.

```

root@master:~# kubectl describe svc backend-service
Name: backend-service
Namespace: default
Labels: app=backend-deploy
Annotations: <none>
Selector: app=backend-deploy
Type: ClusterIP
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.109.215.53
IPs: 10.109.215.53
Port: <unset> 80/TCP
TargetPort: 80/TCP
Endpoints: 192.168.189.73:80,192.168.189.82:80,192.168.235.132:80
Session Affinity: None
Events: <none>
root@master:~#
root@master:~# kubectl get pods -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
backend-deploy-5489dcddc6-48c74 1/1 Running 0 21m 192.168.189.82 worker2 <none> <none>
backend-deploy-5489dcddc6-dggnx 1/1 Running 0 21m 192.168.235.132 worker1 <none> <none>
backend-deploy-5489dcddc6-jvdpj 1/1 Running 0 21m 192.168.189.73 worker2 <none> <none>
root@master:~#

```

So our ClusterIP service has been created and it is targeting the pods which are having the label **app=backend-deploy**. Under the endpoints field we can see the targeted pods.

We can access the pods using the service IP address within the cluster.

```
Curl 10.109.215.53
```

```

root@master:~# curl 10.109.215.53
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@master:~#

```

Now delete the service using below command.

```
kubectl delete svc backend-service
```


3. Creating a deployment and exposing it to a LoadBalancer service

It exposes the Service externally using a cloud provider's load balancer. It is same as NodePort but instead of using the node's IP address, here we will get a single DNS or IP address which cloud provider create for us. We use this service if we are going for PAAS services like EKS or GKE etc.

Step1: Create a deployment

Use the below command to create a deployment.

```
kubectl create deploy frontend-deploy --image nginx --replicas=2
```

```
root@ip-172-31-33-180:~#
root@ip-172-31-33-180:~# kubectl describe deploy frontend-deploy
Name: frontend-deploy
Namespace: default
CreationTimestamp: Tue, 17 Jan 2023 12:36:28 +0000
Labels: app=frontend-deploy
Annotations: deployment.kubernetes.io/revision: 1
Selector: app=frontend-deploy
Replicas: 2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels: app=frontend-deploy
  Containers:
    nginx:
      Image: nginx
      Port: <none>
      Host Port: <none>
      Environment: <none>
      Mounts: <none>
      Volumes: <none>
  Conditions:
    Type Status Reason
    ----
    Available True MinimumReplicasAvailable
    Progressing True NewReplicaSetAvailable
  OldReplicaSets: <none>
  NewReplicaSet: frontend-deploy-7fc64947c6 (2/2 replicas created)
Events:
  Type Reason Age From Message
  ----
  Normal ScalingReplicaSet 44s deployment-controller Scaled up replica set frontend-deploy-7fc64947c6 to 2
```

The above output shows us the pods will be using the label **app=frontend-deploy**.

Step2: Create a LoadBalancer service

We will be using the below command or a Yaml file to create a LoadBalancer service. Cloud provider will automatically create a physical Load balancer for us.

```
kubectl expose deploy frontend-deploy --name frontend-service --type
LoadBalancer --port 80
```

Or create a template from the above command.

```
kubectl expose deploy frontend-deploy --name frontend-service --type
LoadBalancer --port 80 --dry-run=client -o yaml > frontend-service.yaml
```

```

apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: frontend-deploy
  name: frontend-service
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: frontend-deploy
  type: LoadBalancer
status:
  loadBalancer: {}

```

Use the below command to apply the definition file.

```
kubectl apply -f frontend-service.yaml
```

We can see below that a Load balancer service has been created and AWS create a physical load balancer.

```

root@ip-172-31-33-180:~#
root@ip-172-31-33-180:~# kubectl expose deploy frontend-deploy --name frontend-service --type LoadBalancer --port 80
service/frontend-service exposed
root@ip-172-31-33-180:~#
root@ip-172-31-33-180:~# kubectl get svc

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend-service	LoadBalancer	10.100.17.132	abebdc997c05846229a39eb58b33c916-1411797526.us-east-2.elb.amazonaws.com	80:32388/TCP	12s
kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	79m

```

root@ip-172-31-33-180:~#

```

```

root@ip-172-31-33-180:~# kubectl describe svc frontend-service
Name: frontend-service
Namespace: default
Labels: app=frontend-deploy
Annotations: <none>
Selector: app=frontend-deploy
Type: LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.100.17.132
IPs: 10.100.17.132
LoadBalancer Ingress: abebdc997c05846229a39eb58b33c916-1411797526.us-east-2.elb.amazonaws.com
Port: <unset> 80/TCP
TargetPort: 80/TCP
NodePort: <unset> 32388/TCP
Endpoints: 172.31.2.150:80,172.31.22.198:80
Session Affinity: None
External Traffic Policy: Cluster
Events:
  Type    Reason              Age   From             Message
  ----    -
  Normal  EnsuringLoadBalancer 52s   service-controller Ensuring load balancer
  Normal  EnsuredLoadBalancer 49s   service-controller Ensured load balancer
root@ip-172-31-33-180:~#

```