

Helm

Introduction:

Helm is an open-source CNCF project which is written in GO language and originally known as *package manager for Kubernetes*. It is focused on automating the Kubernetes applications lifecycle in a simple and consistent way.

The main objective of Helm as package manager is to manage the Kubernetes resources in automated way. We can install, update, or uninstall the packages for Kubernetes applications, and deploy them with just a few commands.

Writing and maintaining Kubernetes YAML manifest for all the required Kubernetes objects can be a time consuming and tedious task for the simplest of deployments. helm simplifies the process and create a single package that can be advertised to your cluster.

Helm key components:

Chart: It is simply Kubernetes YAML manifest combined into a single package that can be advertised to the cluster. A chart is a helm package. it contains all of the resource definitions necessary to run an application, tool or service inside of a cluster. like yum, apt etc.

Release: A release is an instance of a chart running in the cluster. One chart can often be installed multiple times into the same cluster and each time it is installed, a new release is created.

Repository: It is a location where packaged charts can be stored and shared.

Objectives:

1. Helm Installation
2. Create a Helm Chart
3. Install a Helm Chart
4. Upgrade the Helm chart
5. Rollback to earlier revision
6. Delete the Helm chart
7. Create a Helm repository
8. Deploy a WordPress application from repo
9. Deploy WordPress application using external Database
10. Helm plugins

1. Helm Installation:

We are going to install Helm on our machine. Use the below script to install Helm.

```
curl -fsSL -o get_helm.sh
https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3
chmod 700 get_helm.sh
./get_helm.sh
```

Once it is installed, you can check the version using below command.

```
helm version
```

```
root@master:~#
root@master:~# curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3
root@master:~# chmod 700 get_helm.sh
root@master:~# ./get_helm.sh
Downloading https://get.helm.sh/helm-v3.11.0-linux-amd64.tar.gz
Verifying checksum... Done.
Preparing to install helm into /usr/local/bin
helm installed into /usr/local/bin/helm

root@master:~#
root@master:~# helm version
version.BuildInfo{Version:"v3.11.0", GitCommit:"472c5736ab01133de504a826bd9ee12cbe4e7904", GitTreeState:"clean", GoVersion:"go1.18.10"}
root@master:~#
```

The above output shows that Helm has been installed in our cluster.

2. Create a Helm Chart:

Now let's create our first Helm chart using the below command.

```
helm create project1
```

Above command will create a **project1** directory which will include the files which can be seen in below output.

```

root@master:~/helm# helm create project1
Creating project1
root@master:~/helm#
root@master:~/helm# ls
project1
root@master:~/helm#
root@master:~/helm# tree project1/
project1/
├── Chart.yaml
├── charts
├── templates
│   ├── NOTES.txt
│   ├── _helpers.tpl
│   ├── deployment.yaml
│   ├── hpa.yaml
│   ├── ingress.yaml
│   ├── service.yaml
│   ├── serviceaccount.yaml
│   └── tests
│       └── test-connection.yaml
└── values.yaml

3 directories, 10 files
root@master:~/helm#

```

Under that project directory, we have 3 directories and 10 files.

Chart.yaml file has the information about the project.

```

root@master:~/helm/project1# cat Chart.yaml
apiVersion: v2
name: project1
description: A Helm chart for Kubernetes

# A chart can be either an 'application' or a 'library' chart.
#
# Application charts are a collection of templates that can be packaged into versioned archives
# to be deployed.
#
# Library charts provide useful utilities or functions for the chart developer. They're included as
# a dependency of application charts to inject those utilities and functions into the rendering
# pipeline. Library charts do not define any templates and therefore cannot be deployed.
type: application

# This is the chart version. This version number should be incremented each time you make changes
# to the chart and its templates, including the app version.
# Versions are expected to follow Semantic Versioning (https://semver.org/)
version: 0.1.0

# This is the version number of the application being deployed. This version number should be
# incremented each time you make changes to the application. Versions are not expected to
# follow Semantic Versioning. They should reflect the version the application is using.
# It is recommended to use it with quotes.
appVersion: "1.16.0"

```

Now let's explore the main file which is **values.yaml**.

```
# Default values for project1.

# This is a YAML-formatted file.

# Declare variables to be passed into your templates.

replicaCount: 1

image:

  repository: nginx

  pullPolicy: IfNotPresent

  # Overrides the image tag whose default is the chart appVersion.

  tag: ""

imagePullSecrets: []

nameOverride: ""

fullnameOverride: ""

serviceAccount:

  # Specifies whether a service account should be created

  create: true

  # Annotations to add to the service account

  annotations: {}

  # The name of the service account to use.

  # If not set and create is true, a name is generated using the fullname template

  name: ""

podAnnotations: {}

podSecurityContext: {}

  # fsGroup: 2000

securityContext: {}
```

```
# capabilities:

# drop:

# - ALL

# readOnlyRootFilesystem: true

# runAsNonRoot: true

# runAsUser: 1000

service:

  type: ClusterIP

  port: 80

ingress:

  enabled: false

  className: ""

  annotations: {}

    # kubernetes.io/ingress.class: nginx

    # kubernetes.io/tls-acme: "true"

  hosts:

    - host: chart-example.local

      paths:

        - path: /

          pathType: ImplementationSpecific

  tls: []

  # - secretName: chart-example-tls

  # hosts:

  #   - chart-example.local
```

```
resources: {}

# We usually recommend not to specify default resources and to leave this as a conscious
# choice for the user. This also increases chances charts run on environments with little
# resources, such as Minikube. If you do want to specify resources, uncomment the
# following
# lines, adjust them as necessary, and remove the curly braces after 'resources:'.

# limits:

#   cpu: 100m

#   memory: 128Mi

# requests:

#   cpu: 100m

#   memory: 128Mi

autoscaling:

  enabled: false

  minReplicas: 1

  maxReplicas: 100

  targetCPUUtilizationPercentage: 80

  # targetMemoryUtilizationPercentage: 80

nodeSelector: {}

tolerations: []

affinity: {}
```

The values.yaml file is the main file, where we are going to make the changes. So instead of changing the values in each file, we will change in values.yaml file. Once we are done with the changes, we can install this and all the changes will be reflected automatically.

You can see that in values.yaml, we have image, replicaCount, serviceaccount, service, securityContext etc which are required by templates.

Under templates directory, we have below files which will fetch the values from values.yaml file.

```
root@master:~/helm/project1#  
root@master:~/helm/project1# tree templates/  
templates/  
├── NOTES.txt  
├── _helpers.tpl  
├── deployment.yaml  
├── hpa.yaml  
├── ingress.yaml  
├── service.yaml  
├── serviceaccount.yaml  
├── tests  
│   └── test-connection.yaml  
└──  
  
1 directory, 8 files  
root@master:~/helm/project1#
```

Let's explore deployment.yaml from above templates.

```
apiVersion: apps/v1  
  
kind: Deployment  
  
metadata:  
  name: {{ include "project1.fullname" . }}  
  labels:  
    {{- include "project1.labels" . | nindent 4 }}  
  
spec:  
  {{- if not .Values.autoscaling.enabled }}  
  replicas: {{ .Values.replicaCount }}  
  {{- end }}  
  
selector:
```

```
matchLabels:

  {{- include "project1.selectorLabels" . | nindent 6 }}

template:

  metadata:

    {{- with .Values.podAnnotations }}

    annotations:

      {{- toYaml . | nindent 8 }}

    {{- end }}

    labels:

      {{- include "project1.selectorLabels" . | nindent 8 }}

  spec:

    {{- with .Values.imagePullSecrets }}

    imagePullSecrets:

      {{- toYaml . | nindent 8 }}

    {{- end }}

    serviceAccountName: {{ include "project1.serviceAccountName" . }}

    securityContext:

      {{- toYaml .Values.podSecurityContext | nindent 8 }}

    containers:

      - name: {{ .Chart.Name }}

        securityContext:

          {{- toYaml .Values.securityContext | nindent 12 }}

        image: "{{ .Values.image.repository }}:{{ .Values.image.tag | default
.Chart.AppVersion }}"
```



```
imagePullPolicy: {{ .Values.image.pullPolicy }}

ports:

  - name: http

    containerPort: {{ .Values.service.port }}

    protocol: TCP

livenessProbe:

  httpGet:

    path: /

    port: http

readinessProbe:

  httpGet:

    path: /

    port: http

resources:

  {{- toYaml .Values.resources | nindent 12 }}

{{- with .Values.nodeSelector }}

nodeSelector:

  {{- toYaml . | nindent 8 }}

{{- end }}

{{- with .Values.affinity }}

affinity:

  {{- toYaml . | nindent 8 }}

{{- end }}

{{- with .Values.tolerations }}
```

```
tolerations:

  {{- toYaml . | nindent 8 }}

{{- end }}
```

So this deployment file will fetch the details like image, replicaCount, securityContext etc from values.yaml.

Service.yaml

```
apiVersion: v1

kind: Service

metadata:

  name: {{ include "project1.fullname" . }}

  labels:

    {{- include "project1.labels" . | nindent 4 }}

spec:

  type: {{ .Values.service.type }}

  ports:

    - port: {{ .Values.service.port }}

      targetPort: http

      protocol: TCP

      name: http

  selector:

    {{- include "project1.selectorLabels" . | nindent 4 }}
```

This **service.yaml** file will fetch **service type**, **service port** etc details from **values.yaml**

Similarly, other files will get the details from values.yaml.

3. Install a Helm Chart:

We have our chart with us, let's install this chart using the below command.

We can use lint command to check if our chart is having any errors.

```
helm lint project1
```

```
root@master:~/helm#  
root@master:~/helm# helm lint project1  
==> Linting project1  
[INFO] Chart.yaml: icon is recommended  
  
1 chart(s) linted, 0 chart(s) failed  
root@master:~/helm#
```

Use the below command to check what all resources will be installed. It validates the YAMLs file without connecting with Kube-apiserver.

```
helm template project1
```

Use the below dry-run command to connect with kube-apiserver but it will not install the resources.

```
helm install first-project --debug --dry-run project1
```

Now let's install our first release using the below command.

```
helm install first-project project1
```

```
root@master:~/helm#  
root@master:~/helm# helm install first-project project1  
NAME: first-project  
LAST DEPLOYED: Fri Jan 20 10:22:56 2023  
NAMESPACE: default  
STATUS: deployed  
REVISION: 1  
NOTES:  
1. Get the application URL by running these commands:  
  export POD_NAME=$(kubectl get pods --namespace default -l "app.kubernetes.io/name=project1,app.kubernetes.io/instance=first-project" -o jsonpath="{.items[0].metadata.name}")  
  export CONTAINER_PORT=$(kubectl get pod --namespace default $POD_NAME -o jsonpath="{.spec.containers[0].ports[0].containerPort}")  
  echo "Visit http://127.0.0.1:8080 to use your application"  
  kubectl --namespace default port-forward $POD_NAME 8080:$CONTAINER_PORT  
root@master:~/helm#
```

```
root@master:~/helm# helm list  
NAME          NAMESPACE   REVISION   UPDATED                               STATUS   CHART           APP VERSION  
first-project default     1          2023-01-20 10:22:56.75409116 +0000 UTC deployed project1-0.1.0 1.16.0  
root@master:~/helm#
```

Our first chart is ready with the information present in values.yaml.

Check the helm chart installed using the below command.

```
helm list
```

So our first-project is ready with revision 1.

```

root@master:~/helm#
root@master:~/helm# kubectl get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/first-project-project1-88578c9b-rdw8c   1/1     Running   0           4m25s

NAME                                TYPE               CLUSTER-IP      EXTERNAL-IP   PORT(S)    AGE
service/first-project-project1         ClusterIP          10.97.112.205   <none>        80/TCP     4m25s
service/kubernetes                     ClusterIP          10.96.0.1       <none>        443/TCP    3d23h

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/first-project-project1  1/1     1             1           4m25s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/first-project-project1-88578c9b  1         1         1       4m25s
root@master:~/helm#

```

Above we can see that now we have one deployment with one pod, a replica set and a service.

4. Upgrade the Helm chart

Now let's make some changes in the **values.yaml** file and install the project again.

```

# Default values for project1.
# This is a YAML-formatted file.
# Declare variables to be passed into your templates.

replicaCount: 2

image:
  repository: nginx
  pullPolicy: IfNotPresent
  # Overrides the image tag whose default is the chart appVersion.
  tag: ""

imagePullSecrets: []
nameOverride: ""
fullnameOverride: ""

serviceAccount:
  # Specifies whether a service account should be created
  create: true
  # Annotations to add to the service account
  annotations: {}
  # The name of the service account to use.
  # If not set and create is true, a name is generated using the fullname template
  name: ""

```

We have changed the replicaCount from 1 to 2. Now upgrade the first-project with this latest information. Use the below command.

```
helm upgrade first-project project1
```

It will create another revision which can be seen in below output.

```

root@master:~/helm#
root@master:~/helm# helm upgrade first-project project1
Release "first-project" has been upgraded. Happy Helming!
NAME: first-project
LAST DEPLOYED: Fri Jan 20 10:31:42 2023
NAMESPACE: default
STATUS: deployed
REVISION: 2
NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace default -l "app.kubernetes.io/name=project1,app.kubernetes.io/instance=first-project" -o jsonpath="{.items[0].metadata.name}")
  export CONTAINER_PORT=$(kubectl get pod --namespace default $POD_NAME -o jsonpath="{.spec.containers[0].ports[0].containerPort}")
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl --namespace default port-forward $POD_NAME 8080:$CONTAINER_PORT
root@master:~/helm#

```

```

root@master:~/helm#
root@master:~/helm# helm list
NAME                NAMESPACE    REVISION    UPDATED                               STATUS          CHART           APP VERSION
first-project       default       2           2023-01-20 10:31:42.738567088 +0000 UTC deployed       project1-0.1.0  1.16.0
root@master:~/helm#

```

```

root@master:~/helm#
root@master:~/helm# kubectl get all
NAME                                                    READY   STATUS    RESTARTS   AGE
pod/first-project-project1-88578c9b-hvdkf             1/1     Running   0          4m29s
pod/first-project-project1-88578c9b-rdw8c             1/1     Running   0          13m

NAME                TYPE          CLUSTER-IP    EXTERNAL-IP   PORT(S)    AGE
service/first-project-project1  ClusterIP     10.97.112.205 <none>        80/TCP     13m
service/kubernetes           ClusterIP     10.96.0.1     <none>        443/TCP    3d23h

NAME                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/first-project-project1  2/2     2            2          13m

NAME                DESIRED   CURRENT   READY   AGE
replicaset.apps/first-project-project1-88578c9b  2         2         2       13m
root@master:~/helm#

```

Above we can see that now we have 2 pods as per the information provided to values.yaml.

5. Rollback to earlier revision:

We can use the below command to go back to earlier version, however it will create another revision and rollback the changes.

```
helm rollback first-project 1
```

```

root@master:~/helm#
root@master:~/helm# helm list -a
NAME                NAMESPACE    REVISION    UPDATED                               STATUS          CHART           APP VERSION
first-project       default       2           2023-01-20 10:31:42.738567088 +0000 UTC deployed       project1-0.1.0  1.16.0
root@master:~/helm#
root@master:~/helm#
root@master:~/helm# helm rollback first-project 1
Rollback was a success! Happy Helming!
root@master:~/helm#
root@master:~/helm#
root@master:~/helm# helm list -a
NAME                NAMESPACE    REVISION    UPDATED                               STATUS          CHART           APP VERSION
first-project       default       3           2023-01-20 10:41:25.52743425 +0000 UTC deployed       project1-0.1.0  1.16.0
root@master:~/helm#
root@master:~/helm# kubectl get all
NAME                                                    READY   STATUS    RESTARTS   AGE
pod/first-project-project1-88578c9b-rdw8c             1/1     Running   0          18m

NAME                TYPE          CLUSTER-IP    EXTERNAL-IP   PORT(S)    AGE
service/first-project-project1  ClusterIP     10.97.112.205 <none>        80/TCP     18m
service/kubernetes           ClusterIP     10.96.0.1     <none>        443/TCP    3d23h

NAME                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/first-project-project1  1/1     1            1          18m

NAME                DESIRED   CURRENT   READY   AGE
replicaset.apps/first-project-project1-88578c9b  1         1         1       18m
root@master:~/helm#

```

So, above output shows that we are back to one pod.

6. Delete the Helm chart:

Use the below command to delete the installed Helm chart.

```
helm delete first-project
```

```
root@master:~/helm#  
root@master:~/helm# helm delete first-project  
release "first-project" uninstalled  
root@master:~/helm#  
root@master:~/helm#  
root@master:~/helm#  
root@master:~/helm# helm list  
NAME      NAMESPACE      REVISION      UPDATED STATUS  CHART          APP VERSION  
root@master:~/helm#  
root@master:~/helm# kubectl get pods  
No resources found in default namespace.  
root@master:~/helm#  
root@master:~/helm#
```

Now we do not have any pod, deployment and service in our cluster, as the helm chart has been deleted.

7. Create a Helm repository:

As of now, we do not have any repository added in our cluster. You can check that using below command.

```
helm repo list
```

Search for the repository from the hub, use the below command.

```
helm search hub
```

It will show us all the repositories in our hub.

```
helm search hub wordpress --max-col-width=0
```

Above command will search and list the repo which has the WordPress chart.

Now let's add the Bitnami repository in our cluster and name it **bitnami-repo**.

```
helm repo add bitnami-repo https://charts.bitnami.com/bitnami
```

```

root@master:~/helm#
root@master:~/helm# helm repo add bitnami-repo https://charts.bitnami.com/bitnami
"bitnami-repo" has been added to your repositories
root@master:~/helm#
root@master:~/helm# helm repo list
NAME          URL
bitnami-repo  https://charts.bitnami.com/bitnami
root@master:~/helm#

```

Our repo has been added.

8. Deploy a WordPress application from repo

Let's search for the WordPress chart from the repo

```
helm search repo wordpress
```

Above command will show us the version of Wordpress chart available in our repo.

```

root@master:~/helm# helm search repo wordpress
NAME                CHART VERSION  APP VERSION  DESCRIPTION
bitnami-repo/wordpress  15.2.31        6.1.1        WordPress is the world's most popular blogging ...
bitnami-repo/wordpress-intel  2.1.31        6.1.1        DEPRECATED WordPress for Intel is the most popu...
root@master:~/helm#

```

Use the below command to get the instructions to install this package.

```
helm show readme bitnami-repo/wordpress --version 15.2.31
```

Now let's check the values which will be used in this chart.

```
helm show values bitnami-repo/wordpress --version 15.2.31
```

This application is using the dynamic volumes so we would be creating this application on PAAS service like GKE or EKS.

Use the below command to install Kubernetes.

```

helm install my-release \

--set wordpressUsername=admin \

--set wordpressPassword=password \

--set mariadb.auth.rootPassword=secretpassword \

bitnami-repo/wordpress

```

Above we are providing the values in the command line itself.

```
helm install my-release \
  --set wordpressUsername=admin \
  --set wordpressPassword=password \
  --set mariadb.auth.rootPassword=secretpassword \
  bitnami-repo/wordpress
```

```
W0120 12:36:01.532019 1055 warnings.go:70] Autopilot increased resource requests for Deployment default/my-release-wordpress to meet requirements. See http://g.co/gke/autopilot-resources
W0120 12:36:01.938692 1055 warnings.go:70] Autopilot set default resource requests for StatefulSet default/my-release-mariadb, as resource requests were not specified. See http://g.co/gke/a
NAME: my-release
LAST DEPLOYED: Fri Jan 20 12:35:44 2023
NAMESPACE: default
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
CHART NAME: wordpress
CHART VERSION: 15.2.31
APP VERSION: 6.1.1

** Please be patient while the chart is being deployed **

Your WordPress site can be accessed through the following DNS name from within your cluster:

    my-release-wordpress.default.svc.cluster.local (port 80)

To access your WordPress site from outside the cluster follow the steps below:

1. Get the WordPress URL by running these commands:

    NOTE: It may take a few minutes for the LoadBalancer IP to be available.
    Watch the status with: 'kubectl get svc --namespace default -w my-release-wordpress'

    export SERVICE_IP=$(kubectl get svc --namespace default my-release-wordpress --include "{{ range (index .status.loadBalancer.ingress 0) }}{{ . }}{{ end }}" )
    echo "WordPress URL: http://${SERVICE_IP}/"
    echo "WordPress Admin URL: http://${SERVICE_IP}/admin"

2. Open a browser and access WordPress using the obtained URL.

3. Login with the following credentials below to see your blog:

    echo Username: admin
    echo Password: $(kubectl get secret --namespace default my-release-wordpress -o jsonpath="{.data.wordpress-password}" | base64 -d)
```

Once installed, it will show us the above output.

```
--
$ kubectl get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/my-release-mariadb-0            1/1     Running   0           16m
pod/my-release-wordpress-7c4555f8-dv865  1/1     Running   0           16m

NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)                                AGE
service/kubernetes                  ClusterIP           10.8.128.1      <none>            443/TCP                                26m
service/my-release-mariadb          ClusterIP           10.8.128.106    <none>            3306/TCP                                16m
service/my-release-wordpress        LoadBalancer       10.8.130.25     34.67.91.217     80:31068/TCP,443:30938/TCP            16m

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/my-release-wordpress 1/1     1             1           16m

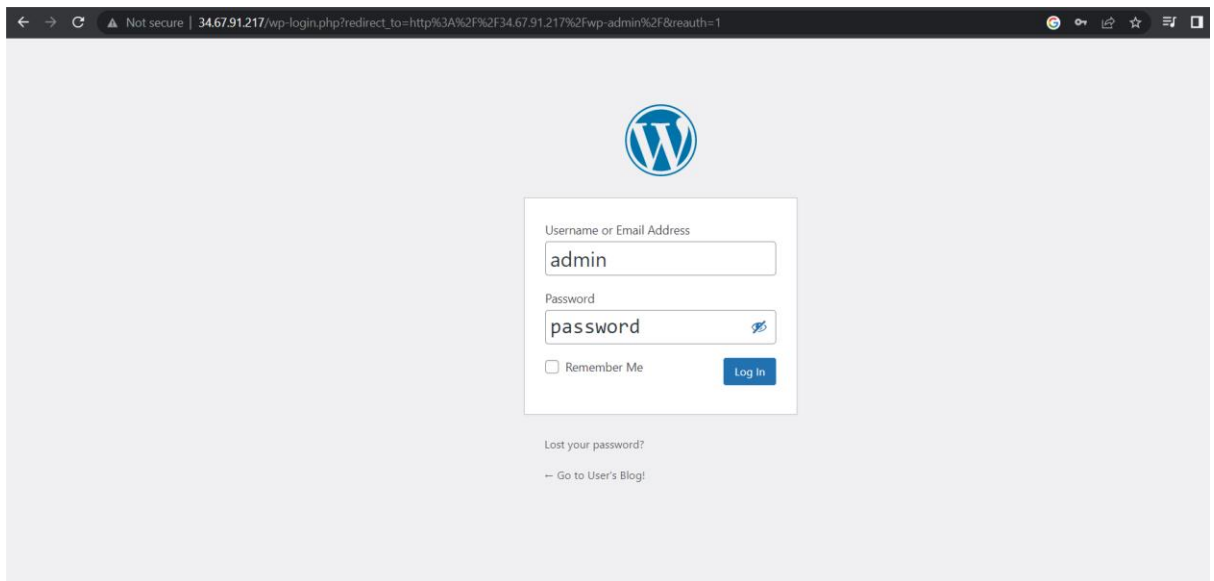
NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/my-release-wordpress-7c4555f8 1         1         1       16m

NAME                                READY   AGE
statefulset.apps/my-release-mariadb 1/1     16m
```

Now our application is ready to be accessed. We can access the application using the service which is exposing our frontend.

Above the service has received the IP address 34.67.91.217, let's open this in our browser using below.

<https://34.67.91.217/admin>



9. Deploy WordPress application using an external Database:

We are going to use an external database and we will provide the values from a file while installing our WordPress application.

Step 1:

Let's install the MySQL database:

Use the below command first to check the chart in our repo.

```
helm search repo mysql
```

```
root@master:~/helm#
root@master:~/helm# helm search repo mysql
NAME                CHART VERSION  APP VERSION  DESCRIPTION
bitnami-repo/mysql  9.4.6          8.0.31       MySQL is a fast, reliable, scalable, and easy t...
bitnami-repo/phpmyadmin  10.4.0        5.2.0        phpMyAdmin is a free software tool written in P...
bitnami-repo/mariadb  11.4.4        10.6.11      MariaDB is an open source, community-developed ...
bitnami-repo/mariadb-galera  7.4.11       10.6.11      MariaDB Galera is a multi-primary database clus...
```

We would get above chart list, now we are going to install mysql and will set some of the values.

```
helm install my-release --set
auth.rootPassword=secretpassword,auth.database=app_database bitnami-
repo/mysql
```

Use the above command to install MySQL. We have provided below values during the installation which will be used by our WordPress application.

```
auth.rootPassword=secretpassword
```

auth.database=app_database

```
$ kubectl get all
NAME                READY   STATUS    RESTARTS   AGE
pod/my-release-mysql-0    1/1     Running   0           2m25s

NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes   ClusterIP     10.8.128.1   <none>         443/TCP    148m
service/my-release-mysql   ClusterIP     10.8.130.152 <none>         3306/TCP   2m26s
service/my-release-mysql-headless ClusterIP      None         <none>         3306/TCP   2m26s

NAME                READY   AGE
statefulset.apps/my-release-mysql  1/1     2m26s
```

Our database is ready now.

Step 2: Install WordPress application

Let's create a file named **wordpress.values.yaml**.

```
vi wordpress.values.yaml
```

And paste the below information.

```
wordpressUsername: techlanders

wordpressPassword: techlanders123

wordpressEmail: contactus@techlanders.com

wordpressFirstName: techlanders

wordpressLastName: training

wordpressBlogName: techlanders.com

service:

  type: LoadBalancer

externalDatabase:

  host: my-release-mysql.default.svc.cluster.local

  user: root

  password: secretpassword

  database: app_database
```

```
mariadb:
```

```
enabled: false
```

Above we have provided credentials for WordPress admin and external database. This helm chart is using mariadb so we are disabling it as we are going for external database.

Let's install the helm chart now.

```
helm install wordpress bitnami-repo/wordpress --values wordpress.values.yaml
```

NAME	READY	STATUS	RESTARTS	AGE
pod/my-release-mysql-0	1/1	Running	0	15m
pod/wordpress-6b48fd45cb-6k5xc	1/1	Running	0	101s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.8.128.1	<none>	443/TCP	161m
service/my-release-mysql	ClusterIP	10.8.130.152	<none>	3306/TCP	15m
service/my-release-mysql-headless	ClusterIP	None	<none>	3306/TCP	15m
service/wordpress	LoadBalancer	10.8.130.255	34.67.91.217	80:32441/TCP,443:30126/TCP	102s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/wordpress	1/1	1	1	102s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/wordpress-6b48fd45cb	1	1	1	103s

NAME	READY	AGE
statefulset.apps/my-release-mysql	1/1	15m

Now our application is ready and we will access it using the service.

In our case, we are having the IP address **34.67.91.217/admin** to access the WordPress application.

10. Helm Plugins:

We can use Helm plugins to understand more about Helm release. In this exercise, we are going to use diff plugin which is used to show the difference between two versions.

To get the list of plugins installed, we can use below command.

```
helm plugin list
```

Let's install diff plugin, use below command.

```
helm plugin install https://github.com/databus23/helm-diff
```

```

root@master:~#
root@master:~# helm plugin list
NAME      VERSION DESCRIPTION
root@master:~#

root@master:~# helm plugin install https://github.com/databus23/helm-diff
Downloading https://github.com/databus23/helm-diff/releases/latest/download/helm-diff-linux-amd64.tgz
Preparing to install into /root/.local/share/helm/plugins/helm-diff
helm-diff installed into /root/.local/share/helm/plugins/helm-diff/helm-diff

The Helm Diff Plugin

* Shows a diff explaining what a helm upgrade would change:
  This fetches the currently deployed version of a release
  and compares it to a local chart plus values. This can be
  used visualize what changes a helm upgrade will perform.

* Shows a diff explaining what had changed between two revisions:
  This fetches previously deployed versions of a release
  and compares them. This can be used visualize what changes
  were made during revision change.

* Shows a diff explaining what a helm rollback would change:
  This fetches the currently deployed version of a release
  and compares it to the previously deployed version of the release, that you
  want to rollback. This can be used visualize what changes a
  helm rollback will perform.

Usage:
  diff [flags]
  diff [command]

Available Commands:
  completion  Generate the autocompletion script for the specified shell
  release     Shows diff between release's manifests
  revision    Shows diff between revision's manifests
  rollback    Show a diff explaining what a helm rollback could perform
  upgrade     Show a diff explaining what a helm upgrade would change.
  version     Show version of the helm diff plugin

```

```

root@master:~#
root@master:~# helm plugin list
NAME      VERSION DESCRIPTION
diff      3.6.0    Preview helm upgrade changes as a diff
root@master:~#

```

Now our plugin has been installed.

Let's explore this by creating our first project using helm chart.

Use the below command.

```
helm install first-project project
```

Below output shows us that we have a deployment created with 2 pods which are exposed using a ClusterIP service.

```

root@master:~/helm#
root@master:~/helm#
root@master:~/helm# helm install first-project project1
NAME: first-project
LAST DEPLOYED: Mon Jan 23 05:22:56 2023
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace default -l "app.kubernetes.io/name=project1,app.kubernetes.io/instance=first-project" -o jsonpath="{.items[0].metadata.name}")
  export CONTAINER_PORT=$(kubectl get pod --namespace default $POD_NAME -o jsonpath="{.spec.containers[0].ports[0].containerPort}")
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl --namespace default port-forward $POD_NAME 8080:$CONTAINER_PORT
root@master:~/helm#

root@master:~/helm#
root@master:~/helm# kubectl get all
NAME                                     READY   STATUS    RESTARTS   AGE
pod/first-project-project1-88578c9b-tpbrm 1/1     Running   0           6s
pod/first-project-project1-88578c9b-wbcm9 1/1     Running   0           6s

NAME                                TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)    AGE
service/first-project-project1      ClusterIP     10.104.221.76   <none>       80/TCP     6s
service/kubernetes                  ClusterIP     10.96.0.1       <none>       443/TCP    6d18h

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/first-project-project1 2/2      2             2           6s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/first-project-project1-88578c9b 2         2         2       6s

```

Let's make changes in values.yaml file.

We have changed the number of replicas and the service from ClusterIP to NodePort, as shown in below picture.

```

replicaCount: 4

image:
  repository: nginx
  pullPolicy: IfNotPresent
  # Overrides the image tag whose default is the chart appVersion.
  tag: ""

imagePullSecrets: []
nameOverride: ""
fullnameOverride: ""

serviceAccount:
  # Specifies whether a service account should be created
  create: true
  # Annotations to add to the service account
  annotations: {}
  # The name of the service account to use.
  # If not set and create is true, a name is generated using the fullname template
  name: ""

podAnnotations: {}

podSecurityContext: {}
  # fsGroup: 2000

securityContext: {}
  # capabilities:
  #   drop:
  #     - ALL
  # readOnlyRootFilesystem: true
  # runAsNonRoot: true
  # runAsUser: 1000

service:
  type: NodePort
  port: 80

```

Let's upgrade our project using below command.

```
Helm upgrade first-project project1
```

Below you can see that our upgrade has been completed.

```
root@master:~/helm#
root@master:~/helm# helm upgrade first-project project1
Release "first-project" has been upgraded. Happy Helming!
NAME: first-project
LAST DEPLOYED: Mon Jan 23 05:26:47 2023
NAMESPACE: default
STATUS: deployed
REVISION: 2
NOTES:
1. Get the application URL by running these commands:
  export NODE_PORT=$(kubectl get --namespace default -o jsonpath="{.spec.ports[0].nodePort}" services first-project-project1)
  export NODE_IP=$(kubectl get nodes --namespace default -o jsonpath="{.items[0].status.addresses[0].address}")
  echo http://$NODE_IP:$NODE_PORT
root@master:~/helm#
```

Now we have a deployment with 4 pods and service has been changed to NodePort.

Our revision has also changed from 1 to 2.

```
root@master:~/helm#
root@master:~/helm# helm list
NAME          NAMESPACE    REVISION    UPDATED                               STATUS    CHART          APP VERSION
first-project default      2           2023-01-23 05:26:47.627645478 +0000 UTC deployed  project1-0.1.0 1.16.0
root@master:~/helm#

root@master:~/helm#
root@master:~/helm# kubectl get all
NAME                                     READY   STATUS    RESTARTS   AGE
pod/first-project-project1-88578c9b-rf5d8 1/1     Running   0          15s
pod/first-project-project1-88578c9b-tf8tq 1/1     Running   0          14s
pod/first-project-project1-88578c9b-tpbrm 1/1     Running   0          4m5s
pod/first-project-project1-88578c9b-wwcm9 1/1     Running   0          4m5s

NAME                                     TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
service/first-project-project1          NodePort    10.104.221.76 <none>         80:30358/TCP    4m5s
service/kubernetes                      ClusterIP    10.96.0.1     <none>         443/TCP         6d18h

NAME                                     READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/first-project-project1  4/4     4            4           4m5s

NAME                                     DESIRED   CURRENT   READY   AGE
replicaset.apps/first-project-project1-88578c9b 4         4         4       4m5s
root@master:~/helm#
root@master:~/helm#
```

Now we can use the below diff command to see the difference between the revisions.

```

root@master:~/helm#
root@master:~/helm# helm diff revision first-project 1 2
default, first-project-project1, Deployment (apps) has changed:
# Source: project1/templates/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: first-project-project1
  labels:
    helm.sh/chart: project1-0.1.0
    app.kubernetes.io/name: project1
    app.kubernetes.io/instance: first-project
    app.kubernetes.io/version: "1.16.0"
    app.kubernetes.io/managed-by: Helm
spec:
-   replicas: 2
+   replicas: 4
  selector:
    matchLabels:
      app.kubernetes.io/name: project1
      app.kubernetes.io/instance: first-project
  template:
    metadata:
      labels:
        app.kubernetes.io/name: project1
        app.kubernetes.io/instance: first-project
    spec:
-   type: ClusterIP
+   type: NodePort
  ports:
    - port: 80
      targetPort: http
      protocol: TCP
      name: http

```

Above output will show us the difference.

We can also check the difference between two releases. Let's create another release from the same chart with few changes like replicaCount=2.

```
helm install second-project project1
```

```

root@master:~/helm#
root@master:~/helm# helm install second-project project1
NAME: second-project
LAST DEPLOYED: Mon Jan 23 06:01:00 2023
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
1. Get the application URL by running these commands:
  export NODE_PORT=$(kubectl get --namespace default -o jsonpath="{.spec.ports[0].nodePort}" services second-project-project1)
  export NODE_IP=$(kubectl get nodes --namespace default -o jsonpath="{.items[0].status.addresses[0].address}")
  echo http://$NODE_IP:$NODE_PORT
root@master:~/helm#

```

```

root@master:~/helm#
root@master:~/helm# kubectl get all

```

NAME	READY	STATUS	RESTARTS	AGE
pod/first-project-project1-88578c9b-rf5d8	1/1	Running	0	36m
pod/first-project-project1-88578c9b-tf8tq	1/1	Running	0	36m
pod/first-project-project1-88578c9b-tpbrm	1/1	Running	0	40m
pod/first-project-project1-88578c9b-wwcm9	1/1	Running	0	40m
pod/second-project-project1-99bd8b666-dv4z4	1/1	Running	0	5s
pod/second-project-project1-99bd8b666-wzsv2	1/1	Running	0	5s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/first-project-project1	NodePort	10.104.221.76	<none>	80:30358/TCP	40m
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	6d19h
service/second-project-project1	NodePort	10.96.102.82	<none>	80:31238/TCP	5s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/first-project-project1	4/4	4	4	40m
deployment.apps/second-project-project1	2/2	2	2	5s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/first-project-project1-88578c9b	4	4	4	40m
replicaset.apps/second-project-project1-99bd8b666	2	2	2	5s

```

root@master:~/helm#

```

```

root@master:~/helm#
root@master:~/helm# helm list

```

NAME	NAMESPACE	REVISION	UPDATED	STATUS	CHART	APP VERSION
first-project	default	2	2023-01-23 05:26:47.627645478 +0000 UTC	deployed	project1-0.1.0	1.16.0
second-project	default	1	2023-01-23 06:03:13.81981629 +0000 UTC	deployed	project1-0.1.0	1.16.0

```

root@master:~/helm#
root@master:~/helm#

```

Above output shows us that we have two releases installed with name **first-project** and **second-project**.

Now let's find out the difference between these releases using below command.

```
helm diff release first-project second-project
```

```

root@master:~/helm#
root@master:~/helm# helm diff release first-project second-project
project1/templates/deployment.yaml has changed:
# Source: project1/templates/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
- name: first-project-project1
+ name: second-project-project1
labels:
  helm.sh/chart: project1-0.1.0
  app.kubernetes.io/name: project1
- app.kubernetes.io/instance: first-project
+ app.kubernetes.io/instance: second-project
  app.kubernetes.io/version: "1.16.0"
  app.kubernetes.io/managed-by: Helm
spec:
- replicas: 4
+ replicas: 2
  selector:
    matchLabels:
      app.kubernetes.io/name: project1
- app.kubernetes.io/instance: first-project
+ app.kubernetes.io/instance: second-project
  template:
    metadata:
      labels:
        app.kubernetes.io/name: project1
- app.kubernetes.io/instance: first-project
+ app.kubernetes.io/instance: second-project
    spec:
- serviceAccountName: first-project-project1
+ serviceAccountName: second-project-project1
project1/templates/service.yaml has changed:
# Source: project1/templates/service.yaml
apiVersion: v1
kind: Service
metadata:
- name: first-project-project1
+ name: second-project-project1
labels:
  helm.sh/chart: project1-0.1.0
  app.kubernetes.io/name: project1
- app.kubernetes.io/instance: first-project
+ app.kubernetes.io/instance: second-project
  app.kubernetes.io/version: "1.16.0"
  app.kubernetes.io/managed-by: Helm
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: http
      protocol: TCP
      name: http
  selector:
    app.kubernetes.io/name: project1
- app.kubernetes.io/instance: first-project
+ app.kubernetes.io/instance: second-project
project1/templates/serviceaccount.yaml has changed:
# Source: project1/templates/serviceaccount.yaml
apiVersion: v1
kind: ServiceAccount
metadata:
- name: first-project-project1
+ name: second-project-project1

```

There is another command, we can see the difference what a rollback can do using the below command.


```
helm diff rollback first-project 1
```

We can uninstall the plugin using below command.

```
helm plugin uninstall diff
```

Use the below command to package our release.

```
helm package project1
```

Use below command to pull a chart from our hub.

```
helm pull bitnami-repo/wordpress
```

Use below command to pull a chart from our hub in unzipped format.

```
helm pull bitnami-repo/wordpress --untar
```

See the output below, we have the packages ready which can be pushed or shared with other team.

```
root@master:~/helm#  
root@master:~/helm# ls  
project1  project1-0.1.0.tgz  wordpress  wordpress-15.2.31.tgz  
root@master:~/helm#
```